


## CS 597 - Report 6

Experimented and set up environment on cooley (Argonne's cluster) for commencement of experiments.

For benchmarking purposes, extended single node caffe to include support for multi-node training using MPI.

Implemented AlexNet on Caffe and performed distributed training experiments on CIFAR-10 dataset. The following plots show the test accuracies of the models as the training progresses. –




After a parameter sweep of batch size and learning rates for each configuration of number of workers the following batch sizes and learning rates were converged upon for each number of workers –

0.5\_lr\_128\_batch\_size\_1w  
0.5\_lr\_128\_batch\_size\_2w  
0.4\_lr\_128\_batch\_size\_4w  
0.6\_lr\_64\_batch\_size\_8w

The criteria used to select the best configuration for a given number of workers takes in to consideration both, the time taken to converge and the final loss attained at convergence.


CS 597 - Report 6

An illustration is presented below for 4 workers –




This represents the progression of L2 loss for 4 workers as training proceeds with respect to time-  
The plot below maps the area under the loss curves for each experiment-

Distributed CIFAR-10 Loss Plot Areas (Performance w.r.t time) over parameter sweep




## CS 597 - Report 6



The histogram above visualizes and sorts each experiment according to the areas of the loss curves.

This whole experiment was repeated for each #worker configuration and plots were obtained for 1, 2, 4 & 8 workers over the CIFAR-10 dataset. The best of each were chosen and trained to obtain the following plot –

## CS 597 - Report 6



As shown previously. Now, our objective is to obtain better performance for our library.

The selective compilation and environment setup has concluded on Cooley. The makefile looks like this – all: program

```
program: cppcode.o
    g++ -std=c++11 -o convnet -L/soft/visualization/cuda-8.0.61/lib64 -lcuda \
    -lcudart -lcudnn -lcublas -lcurand main.cpp FCLayerCU.o FCLayerCPP.o \
    ConvLayerCU.o ConvLayerCPP.o
    rm *.o
```

```
cppcode.o: cudacode.o
    g++ -c -std=c++11 -o FCLayerCPP.o FCLayer.cpp
    g++ -c -std=c++11 -o ConvLayerCPP.o ConvLayer.cpp
```

```
cudacode.o:
    nvcc -c -std=c++11 -arch=sm_37 -o ConvLayerCU.o ConvLayer.cu
    nvcc -c -std=c++11 -arch=sm_37 -o FCLayerCU.o FCLayer.cu
```

```
clean:
    rm -rf *.o convnet
```

However, it was observed that a simple hardware change has resulted in a couple of cuDNN primitives failing to execute resulting in NaN outputs which are currently being tackled.